

# Claude Code Web Pipeline

A custom web research pipeline that replaces Claude Code's built-in WebFetch with clean markdown extraction, project-level caching, and AI-powered triage — keeping full-resolution web content available while consuming minimal context window.

## The Problem




Figure 1: The problem: dumping entire pages into context

Claude Code's built-in web tools either dump entire pages into the conversation (wasteful — a typical page is 20-70K chars) or return lossy summaries (insufficient for detailed work). Neither approach scales when you need to research across multiple pages.

## The Solution

### Step-by-Step

#### 1. Discover — `WebSearch`

Claude Code's built-in web search returns concise summaries and source URLs. Good for discovery, bad for extraction.

#### 2. Fetch + Cache — `batch_read_urls`

A single MCP tool call fetches all discovered URLs in parallel through Jina Reader (`r.jina.ai`). Each page is:

- Rendered through Jina's headless browser
- Converted to clean markdown
- Stripped of junk (nav, cookies, ads, sidebars, modals — 30+ CSS selectors)
- Images removed (markdown-only)
- Cached permanently to `{project}/.web_cache/{hash}.md`

The tool returns **metadata only** — never the content:

```
[{"path": "/project/.web_cache/a1b2c3.md", "title": "Page Title", "lines": 340, "chars": 18200, "cached": true}, {"path": "/project/.web_cache/d4e5f6.md", "title": "Another Page", "lines": 520, "chars": 31000, "cached": false}]
```

#### 3. Triage — Parallel Haiku Agents

Cheap, fast AI agents (one per cached page) scan the full content and return structured JSON identifying exactly which line ranges are relevant to the query:

```
{"relevant": true, "ranges": [[10, 45], [80, 92]]} {"relevant": false}
```

Agents run in parallel — triage takes ~3-5s regardless of page count.

#### 4. Inject — `Read with offset/limit`

Only the relevant line ranges enter the main conversation. Everything else stays on disk, available for later queries.




Figure 2: End-to-end pipeline: WebSearch → batch fetch → cache → haiku triage → selective read

## What Gets Stripped

30+ CSS selectors strip common junk patterns before caching:

Category	Selectors
Navigation	nav, header, footer, [role='banner'], [role='navigation']
Consent	[class*='cookie'], [class*='consent'], [id*='cookie']
Popups	[class*='popup'], [class*='modal'], [class*='overlay']
Marketing	[class*='newsletter'], [class*='subscribe'], [class*='signup']
Ads/Promos	[class*='advert'], [class*='sponsor'], [class*='promo']
Social	[class*='share'], [class*='social'], [class*='comment']
Layout junk	[class*='sidebar'], [class*='widget'], [class*='related-post']

## Cache Format

Files stored as `.web_cache/{sha256[:12]}.md` with YAML frontmatter:

```
---  
url: https://example.com/article  
title: How to Build a Web Pipeline  
fetched: 2026-02-15T10:30:00+00:00  
hash: a1b2c3d4e5f6  
---  
Title: How to Build a Web Pipeline  
URL Source: https://example.com/article  
Markdown Content:  
[... clean article content ...]
```

Cache is **permanent and project-scoped** — pages persist across sessions. Use `list_cache` to see what's available before re-fetching.

## Performance

Scenario	Pages	Cached	Actually Used	Compression	Triage Time
Narrow single-page query	2	24K chars	1.2K chars	95%	~4s
Multi-page comparison	3	29K chars	8.5K chars	70%	~3s
Broad 5-page research	5	83K chars	10.5K chars	87%	~5s

Compression scales with irrelevance — narrow queries save more. Triage latency is flat because agents run in parallel.

## Architecture

Three files. That's it.




Figure 3: CSS selector stripping: raw HTML junk → clean markdown




Figure 4: Context window compression: cached vs actually used

File	Role
<code>~/.claude/mcp-servers/jina-reader.py</code>	The MCP server. 3 tools ( <code>read_url</code> , <code>batch_read_urls</code> , <code>list_cache</code> ), CSS stripping, Jina API integration, caching. Single self-contained Python script — runs via <code>uv run --script</code> with inline PEP 723 deps.
<code>~/.claude.json</code>	Registers the server globally under <code>mcpServers.jina-reader</code> with the <code>JINA_API_KEY</code> env var.
<code>~/.claude/CLAUDE.md</code>	Instructions that tell Claude Code to always use this 5-step workflow for any web task.

## Setup

1. **Get a Jina API key** from [jina.ai](https://jina.ai)
2. **Drop the server script** at `~/.claude/mcp-servers/jina-reader.py`
3. **Register it** in Claude Code:

```
claude mcp add -s user -e "JINA_API_KEY=your_key_here" -- jina-reader \
    uv run --script ~/.claude/mcp-servers/jina-reader.py
```

4. **Add workflow instructions** to `~/.claude/CLAUDE.md`
5. **Restart Claude Code** to pick up the new MCP server

## Known Limitations

- **Cloudflare-protected sites** (Medium, Hacker News) may return challenge pages (~500 chars). Skip these.
- **Nav-heavy sites** (The New Stack) still have junk despite CSS stripping — content may be buried deep.




Figure 5: Architecture: 3 config files → project cache

- **MCP server restarts** are required after editing `jina-reader.py` — restart Claude Code to apply changes.
- **cache\_dir must be absolute** — the MCP server's working directory may not match the project directory.